

Question 2: Rate Limiting and Request Size Limiting with KONG API Gateway

Overview

This guide demonstrates how to implement rate limiting and request size limiting using KONG API Gateway to protect your APIs from abuse and ensure fair resource usage.

Prerequisites

- Docker and Docker Compose installed
- Basic understanding of APIs
- Completed Question 1 (Flask API running)

Step 1: Setup KONG with Docker Compose

Create docker-compose.yml



version: '3.8'

services:

kong-database:

 image: postgres:13

 environment:

 POSTGRES_USER: kong

 POSTGRES_DB: kong

 POSTGRES_PASSWORD: kongpass

 ports:

 - "5432:5432"

 volumes:

 - kong-data:/var/lib/postgresql/data

 networks:

 - kong-net

kong-migration:

 image: kong:3.4

 command: kong migrations bootstrap

 environment:

 KONG_DATABASE: postgres

 KONG_PG_HOST: kong-database

 KONG_PG_USER: kong

 KONG_PG_PASSWORD: kongpass

 depends_on:

 - kong-database

 networks:

 - kong-net

kong:

 image: kong:3.4

 environment:

 KONG_DATABASE: postgres

 KONG_PG_HOST: kong-database

 KONG_PG_USER: kong

 KONG_PG_PASSWORD: kongpass

 KONG_PROXY_ACCESS_LOG: /dev/stdout

 KONG_ADMIN_ACCESS_LOG: /dev/stdout

 KONG_PROXY_ERROR_LOG: /dev/stderr

 KONG_ADMIN_ERROR_LOG: /dev/stderr

 KONG_ADMIN_LISTEN: 0.0.0.0:8001

 ports:

 - "8000:8000" # Proxy port

 - "8443:8443" # Proxy SSL port

```
- "8001:8001" # Admin API port
- "8444:8444" # Admin API SSL port
```

depends_on:

- kong-database
- kong-migration

networks:

- kong-net

volumes:

kong-data:

networks:

kong-net:

driver: bridge

Start KONG



bash

```
# Start all services
```

```
docker-compose up -d
```

```
# Check if KONG is running
```

```
curl -i http://localhost:8001/
```

```
# You should see KONG's admin API response
```

Step 2: Register Your Flask API as a Service



bash

```
# Add your Flask API as a service in KONG
```

```
curl -i -X POST http://localhost:8001/services/ \
--data "name=plagiarism-api" \
--data "url=http://host.docker.internal:5000"
```

```
# Expected response: HTTP/1.1 201 Created with service details
```

Note: Use host.docker.internal to access localhost from Docker container.

Step 3: Create a Route



bash

```
# Create a route to access the service
curl -i -X POST http://localhost:8001/services/plagiarism-api/routes \
--data "name=plagiarism-route" \
--data "paths[]=/api/plagiarism"
```

Now your API is accessible at: http://localhost:8000/api/plagiarism/check

Step 4: Configure Rate Limiting

Option 1: Basic Rate Limiting (Simple)



bash

```
# Add rate limiting plugin (5 requests per minute)
curl -i -X POST http://localhost:8001/services/plagiarism-api/plugins \
--data "name=rate-limiting" \
--data "config.minute=5" \
--data "config.policy=local"
```

This limits each client to 5 requests per minute

Option 2: Advanced Rate Limiting



bash

```
# More granular control
curl -i -X POST http://localhost:8001/services/plagiarism-api/plugins \
--data "name=rate-limiting" \
--data "config.second=2" \
--data "config.minute=10" \
--data "config.hour=100" \
--data "config.policy=local" \
--data "config.fault_tolerant=true"
```

Rate Limits Explained:

- `second=2`: Max 2 requests per second
- `minute=10`: Max 10 requests per minute
- `hour=100`: Max 100 requests per hour
- First limit hit triggers the restriction

Step 5: Configure Request Size Limiting



bash

```
# Limit request body size to 5MB
curl -i -X POST http://localhost:8001/services/plagiarism-api/plugins \
--data "name=request-size-limiting" \
--data "config.allowed_payload_size=5" \
--data "config.size_unit=megabytes"
```

Configuration Options:

- `allowed_payload_size`: Maximum size
- `size_unit`: megabytes, kilobytes, or bytes

Step 6: Test Rate Limiting

Create a test script: `test_rate_limit.py`



python

```

import requests
import time

API_URL = "http://localhost:8000/api/plagiarism/check"

# Test rate limiting
print("Testing Rate Limiting...")
print("==" * 50)

for i in range(15):
    try:
        response = requests.get(API_URL)
        print(f"Request {i+1}: Status {response.status_code}")

        if response.status_code == 429:
            print(f"⚠️ Rate limit exceeded!")
            print(f" Headers: {dict(response.headers)}")
            break

        time.sleep(0.5)
    except Exception as e:
        print(f" Error: {e}")

    print("==" * 50)

```

Expected Output:



Request 1: Status 200
 Request 2: Status 200
 Request 3: Status 200
 Request 4: Status 200
 Request 5: Status 200
 Request 6: Status 429
 ⚠️ Rate limit exceeded!
 Headers: {'X-RateLimit-Limit-Minute': '5', 'X-RateLimit-Remaining-Minute': '0'}

Step 7: Test Request Size Limiting

Create test_size_limit.py



```
import requests
```

```
API_URL = "http://localhost:8000/api/plagiarism/check"
```

```
# Create a large file (6MB - exceeds 5MB limit)
```

```
large_content = "A" * (6 * 1024 * 1024) # 6MB of 'A's
```

```
with open('large_file.txt', 'w') as f:
```

```
    f.write(large_content)
```

```
# Try to upload large file
```

```
print("Testing Request Size Limiting...")
```

```
print("=" * 50)
```

```
try:
```

```
    files = {
```

```
        'original': open('large_file.txt', 'rb'),
```

```
        'submission': open('large_file.txt', 'rb')}
```

```
}
```

```
response = requests.post(API_URL, files=files)
```

```
print(f'Status Code: {response.status_code}')
```

```
if response.status_code == 413:
```

```
    print("✅ Request size limit working!")
```

```
    print(f'Response: {response.text}')
```

```
else:
```

```
    print(f'Response: {response.json()}')
```

```
except Exception as e:
```

```
    print(f'Error: {e}')
```

```
print("=" * 50)
```

Step 8: Monitor and Manage

View All Plugins



```
curl -i http://localhost:8001/plugins
```

View Service Configuration



```
curl -i http://localhost:8001/services/plagiarism-api
```

View Routes



```
curl -i http://localhost:8001/routes
```

Delete a Plugin (if needed)



```
# Get plugin ID first
```

```
curl http://localhost:8001/plugins
```

```
# Delete using ID
```

```
curl -i -X DELETE http://localhost:8001/plugins/{plugin-id}
```

Step 9: Advanced Configuration

Consumer-Based Rate Limiting



```
# Create a consumer
curl -i -X POST http://localhost:8001/consumers \
--data "username=student1"

# Create API key for consumer
curl -i -X POST http://localhost:8001/consumers/student1/key-auth \
--data "key=student1-api-key"

# Enable key authentication
curl -i -X POST http://localhost:8001/services/plagiarism-api/plugins \
--data "name=key-auth"

# Now different rate limits per consumer
curl -i -X POST http://localhost:8001/consumers/student1/plugins \
--data "name=rate-limiting" \
--data "config.minute=20"
```

Screenshots to Capture

Screenshot 1: KONG Admin API Running



```
curl http://localhost:8001/ | jq
```

Screenshot 2: Service Registration



```
curl http://localhost:8001/services/plagiarism-api | jq
```

Screenshot 3: Rate Limiting Plugin Active



```
curl http://localhost:8001/services/plagiarism-api/plugins | jq
```

Screenshot 4: Rate Limit Exceeded (429 Response)

- Run the test_rate_limit.py script
- Capture the 429 response with headers

Screenshot 5: Request Size Limit (413 Response)

- Run the test_size_limit.py script
- Capture the 413 response

Screenshot 6: Successful Request Within Limits

- Show a successful API call with rate limit headers

Verification Checklist

- KONG is running (docker-compose ps)
- Service registered in KONG
- Route created and accessible
- Rate limiting plugin configured
- Request size limiting plugin configured
- Rate limit tested and working (429 response)
- Size limit tested and working (413 response)
- Headers showing rate limit information

Troubleshooting

Issue: Cannot connect to Flask API from KONG

Solution: Use host.docker.internal instead of localhost in service URL

Issue: 404 Not Found

Solution: Check route configuration: curl http://localhost:8001/routes

Issue: Rate limiting not working

Solution:

- Verify plugin is enabled: curl http://localhost:8001/plugins
- Check plugin configuration
- Try different client IPs

Issue: Docker containers not starting

Solution:

- Check logs: docker-compose logs kong
- Ensure ports 8000, 8001, 5432 are not in use
- Run: docker-compose down -v and restart

Cleanup



bash

```
# Stop all services
docker-compose down
```

```
# Remove volumes (careful: deletes all data)
docker-compose down -v
```

```
# Remove images
docker rmi kong:3.4 postgres:13
```

Summary

You have successfully:

1. Set up KONG API Gateway with Docker
2. Registered your Flask API as a service
3. Configured rate limiting (requests per time period)
4. Configured request size limiting (max payload size)
5. Tested both limits with verification scripts
6. Captured evidence screenshots

This implementation protects your API from:

- **Abuse:** Rate limiting prevents excessive requests
- **DoS attacks:** Request size limits prevent memory exhaustion
- **Resource overuse:** Fair usage across all clients